


GOLD PRICE PREDICTION

Rutik Ingale

250840550020

Faculty

Dr. Kiran Waghmare



Problem Statement and Objective

The price of gold changes daily and depends on many financial factors. Predicting these prices manually is difficult and inaccurate. The goal of this project is to develop a machine learning model that can accurately predict future gold prices using historical financial data. The project uses regression techniques to estimate the exact gold price and classification algorithms to determine whether the price will be high or low. The model helps investors and traders make informed decisions.

DataSet Information

1. Dataset Name

Gold Price Dataset (gld_price_data.csv)

2. Source of Dataset

Kaggle.com

3. Number of Records

- Rows: Approximately 2290+
- Columns: 6 main features +

4. Dataset Columns and Their Meaning

1. Date

Type: Object / String

- Represents the trading date for the recorded values.
- Useful for time-series but removed for model training.

2. SPX (S&P 500 Index)

Type: Float

- Tracks performance of top U.S. companies.
- Used because gold price often moves opposite to stock market trends.

3. USO (Crude Oil ETF Prices)

Type: Float

- Represents fluctuations in crude oil prices.
- Gold price is indirectly influenced by oil market trends.

4. SLV (Silver ETF Prices)

Type: Float

- Represents the price of silver.
- Gold and silver prices have strong positive correlation.

5. EUR/USD (Currency Exchange Rate)

Type: Float

- Indicates the strength of US Dollar vs Euro.
- Gold is globally traded in USD; thus currency value impacts gold prices.

6. GLD (Gold ETF Price) → Target for Regression

Type: Float

- Represents the actual price of gold.
- This is the main value you predict using regression.

7. Price_Class (Derived Column) → Target for Classification

Type: Integer (0 or 1)

- Created by splitting GLD price using the median:
 - **1 = High price**
 - **0 = Low price**
- Used for classification models (Logistic Regression, Decision Tree, XGBoost)

5. Data Type Summary

Column Data Type Use

Date	Object	Dropped before training
SPX	Float	Feature
USO	Float	Feature
SLV	Float	Feature
EUR/USD	Float	Feature
GLD	Float	Regression Target
Price_Class	Integer	Classification Target

6. Data Characteristics

- No missing values (checked using isnull().sum())
- Mostly numerical values → suitable for ML models
- Contains financial indicators highly correlated with gold price (checked using correlation heatmap)
- Shows positive correlation between:
 - **GLD ↔ SLV**
 - **GLD ↔ EUR/USD**
- Dataset is clean and ready for ML without additional preprocessing

7. Why This Dataset is Suitable for ML?

Numeric values :- ideal for ML algorithms

High correlation between features :- strong predictive power

Includes multiple financial indicators :- more accurate predictions

Sufficient size for training ML models

Works for both Regression & Classification tasks

Algorithms used and it's justification

A. Regression Algorithm

1. Random Forest Regressor

Why This Algorithm?

- Gold price prediction is a **complex financial problem** with many non-linear relations.
- Random Forest is an **ensemble algorithm** (uses many decision trees).
- It learns patterns from multiple trees and averages them to make a final prediction.

Justification

Handles **non-linear data** extremely well

Reduces **overfitting**

Works effectively with **large datasets**

Gives **high accuracy** compared to single models

Very robust for financial forecasting

Why not Linear Regression?

- Gold data is **not linear**
- Relationship between SPX, oil, silver, currency → GOLD is complex
- Linear regression would give lower accuracy

B. Classification Algorithms

I created a new target variable **Price_Class**

(High = 1, Low = 0).

2. Logistic Regression

Why This Algorithm?

- Best **baseline model** for binary classification.
- Works well when the relationship is somewhat linear.

Justification

Easy to interpret

Fast and efficient

Good starting point for comparison

Helps check if advanced models are needed

3. Decision Tree Classifier

Why This Algorithm?

- Splits data into decision rules.
- Captures **non-linear patterns** very well.

Justification

Understandable visual model (tree plot)

Works well with mixed numerical data

Handles non-linearity better than logistic regression

Can show feature importance

Reason to include

It helps you compare how “simple tree-based models” perform versus advanced boosted models.

4. XGBoost Classifier

Why This Algorithm?

- Most powerful tree-based boosting algorithm.
- Used widely in finance forecasting & Kaggle competitions.
- Boosts weak learners into a strong model.

Justification

High accuracy, often best among all classifiers

Handles noise and outliers well

Fast and optimized model

Learns complex patterns better than Decision Trees

Uses regularization to prevent overfitting

Preprocessing steps

Preprocessing Steps

Preprocessing is an essential stage in the development of any machine learning model. It ensures that the dataset is clean, consistent, and suitable for training predictive models. The following preprocessing steps were applied in this project:

1. Data Loading

The dataset gld_price_data.csv was loaded using the Pandas library. This step converts the raw CSV file into a structured DataFrame for further analysis.

2. Initial Data Inspection

Several commands (head(), tail(), shape, info(), describe()) were used to understand the dataset structure.

This step helped identify:

- The number of rows and columns
- Data types of each column
- Basic statistical properties
- Absence of missing values

3. Removal of Unnecessary Columns

The Date column was removed because it is a non-numeric feature and does not contribute directly to the prediction process. For regression, the GLD column was kept as the target variable. For classification, both Date and GLD were excluded from the feature set.

4. Feature and Target Separation

For Regression:

- Features (X): SPX, USO, SLV, EUR/USD
- Target (y): GLD

For Classification:

A new categorical variable called Price_Class was created by comparing each GLD value with the median price:

- 1 → High Price
- 0 → Low Price

The final classification features included SPX, USO, SLV and EUR/USD, and the target variable was Price_Class.

5. Exploratory Data Analysis (EDA)

To understand the relationship between variables, the following visual analyses were performed:

- Correlation Heatmap: Shows the interdependence between features and the gold price.
- Distribution Plot: Illustrates how the GLD price is distributed across the dataset.

These analyses helped identify important features and detect any skewness or irregular patterns.

6. Train–Test Split

The dataset was divided into:

- 80% Training Data
- 20% Testing Data

This ensures that the model is trained on the majority of the data while still being evaluated on unseen data to check its performance.

The split was done using `train_test_split()` with a fixed random state for reproducibility.

7. Creation of Classification Target Variable

To perform classification, the continuous GLD values were converted into binary categories (high or low).

This step enabled the use of classification algorithms such as Logistic Regression, Decision Tree, and XGBoost.

Code Snippets

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from xgboost import XGBRegressor, XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

# Collection and Processing

gold_data=pd.read_csv('gold_prices_data.csv')

gold_data.head()

   Date      SPX      GLD      USO      SLV  EUR/USD
0 1/2/2008  1447.000034  84.980001  78.470001  15.180  1.471062
1 1/2/2008  1447.000034  85.570000  78.370003  15.285  1.474461
2 1/4/2008  1411.830005  85.120007  77.300008  15.187  1.475492
3 1/7/2008  1418.180054  84.789097  75.500000  15.063  1.488299
4 1/8/2008  1390.139941  88.770999  76.059986  15.590  1.557090

gold_data.tail()

   Date      SPX      GLD      USO      SLV  EUR/USD
2285 5/8/2018  2671.19022  124.589996  14.0000  15.5100  1.180789
2286 5/9/2018  2607.700039  124.330002  14.5700  15.5300  1.194722
2287 5/10/2018  2723.070088  125.180000  14.4100  15.7400  1.191763
2288 5/14/2018  2730.126883  124.489998  14.3800  15.5600  1.193118
2289 5/16/2018  2725.780029  122.543800  14.4058  15.4542  1.182033

gold_data.shape
(2299, 6)

gold_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2299 entries, 0 to 2298
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   Date            2298 non-null    object  
 1   SPX             2298 non-null    float64 
 2   GLD             2298 non-null    float64 
 3   USO             2298 non-null    float64 
 4   SLV             2298 non-null    float64 
 5   EUR/USD         2298 non-null    float64 
dtypes: float64(5), object(1)
memory usage: 107.5+ KB

gold_data.describe()

   SPX      GLD      USO      SLV  EUR/USD
count  2290.000000  2390.000000  2290.000000  2390.000000  2290.000000
mean   1654.316778  122.732875  31.842221  20.084997  1.239553
std    518.111940  23.835485  18.922917  7.002998  0.131947
min    678.530028  70.000000  7.960000  8.550000  0.130407
25%   1238.874098  100.725000  14.380000  15.700000  1.171313
50%   1551.434098  120.580002  33.869998  17.285000  1.302397
75%   2073.010070  132.840004  37.827501  22.825000  1.359071
max    2872.870117  194.589996  117.480003  47.259998  1.598798

correlation = gold_data.drop(['Date'], axis=1).corr()

plt.figure(figsize=(8, 8))
sns.heatmap(correlation, cmap='Blues', annot=True, square=True, fmt=".1f", annot_kws={"size":8})
plt.title("Correlation Heatmap")
plt.show()

Correlation values of gold
print(correlation['GLD'])

SPX      0.002345
GLD      1.000000
USO     -0.186360
SLV      0.006632
```

```


print(correlation['GLD'])

SPX 0.00945
USD -1.00000
EURUSD -0.10000
SV 0.06652
EURUSD -0.054375
Name: GLD, dtype: float64



Checking distribution of gold price


sns.distplot(gold_data['GLD'],color='green')
-- /tmp/ipython-input-1190600815.py:2: UserWarning:
'distplot' is a deprecated function and will be removed in seaborn v0.14.0.
Please update your code to use either 'displot' (a figure-level function with
similar flexibility) or 'histplot' (an axes-level function for histograms).
For a guide to updating your code to use the new functions, please see
https://github.com/mwaskom/seaborn/blob/0d4412e7a2974457ad5172750bae5721
sns.distplot(gold_data['GLD'],color='green')
kws: label='GLD', label=Density'



```

```

Splitting the feature and target
gold_data.drop(['Date','GLD'],axis=1)
y=gold_data['GLD']

print(x)

      SPX    USD    SV    EURUSD
0  1447.160034  78.470001  15.1800  1.474692
1  1447.160034  78.370003  15.2500  1.474691
2  1447.160034  78.270003  15.3200  1.474692
3  1416.188054  75.580000  15.9500  1.468299
4  1398.189041  76.050000  15.9800  1.574999
...
2285 2671.919022  14.860000  15.5100  1.186789
2286 2723.479068  14.410000  15.7700  1.191772
2287 2723.479068  14.410000  15.7400  1.191753
2288 2738.129083  14.380000  15.5600  1.193118
2289 2725.788023  14.460000  15.4500  1.152633

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2)

Model Train

regressor=RandomForestRegressor(n_estimators=100)

#train the model
regressor.fit(x_train,y_train)

RandomForestRegressor(n_estimators=100)

model evaluation

#prediction on test data
test_data_prediction=regressor.predict(x_test)

print(test_data_prediction)

Show hidden output



We coulnd score error.


error_score=metrics.r2_score(y_test,test_data_prediction)
print("R squared error : ",error_score)

R squared error :  0.98913951526177954

y_test=np.sqrt(y_test)



```

```

median_price = gold_data['GLD'].median()
gold_data['Price_Class'] = gold_data['GLD'].apply(lambda x: 1 if x >= median_price else 0)
X = gold_data.drop(['GLD', 'Price_Class', 'Date'], axis=1)
y = gold_data['Price_Class']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

logistic Regression

log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)

logisticRegression(max_iter=1000)
logisticRegression(max_iter=1000)

pred_lr = log_reg.predict(Xc_test)
print(pred_lr)

Show hidden output

acc_lr = accuracy_score(y_test, pred_lr)
print('Logistic Regression Accuracy:', acc_lr)
print(confusion_matrix(y_test, pred_lr))
print(classification_report(y_test, pred_lr))

Logistic Regression Accuracy: 0.925764192139738
[[ 239  219]
 [ 458  458]]
precision recall f1-score support
0    0.93   0.92   0.93    239
1    0.92   0.93   0.92    219
accuracy          0.93   0.93   0.93    458
weighted avg    0.93   0.93   0.93    458

Decision Tree

dtc = DecisionTreeClassifier()
dtc.fit(Xc_train, yc_train)

DecisionTreeClassifier()
DecisionTreeClassifier()

pred_dtc = dtc.predict(Xc_test)
print('Decision Tree Classifier Accuracy:', acc_dtc)
print(confusion_matrix(yc_test, pred_dtc))
print(classification_report(yc_test, pred_dtc))

acc_dtc = accuracy_score(yc_test, pred_dtc)
print('Decision Tree Classifier Accuracy:', acc_dtc)
print(confusion_matrix(yc_test, pred_dtc))
print(classification_report(yc_test, pred_dtc))

Decision Tree Classifier Accuracy: 0.9694323154418484
[[ 239]
 [ 458]]
precision recall f1-score support
0    0.95   0.99   0.97    239
1    0.99   0.95   0.97    219
accuracy          0.97   0.97   0.97    458
weighted avg    0.97   0.97   0.97    458

from sklearn import tree
plt.figure(figsize=(10,10))
tree.plot_tree(dtc,
               filled=True,
               feature_names=Xc.columns,
               class_names=['Low', 'High'],
               rounded=True)
plt.show()

```

```

XG BOOST


xgbc = XGBClassifier()
xgbc.fit(Xc_train, yc_train)

XGBClassifier()
XGBClassifier()

pred_xgbc = xgbc.predict(Xc_test)
print('XGBoost Classifier Accuracy:', acc_xgbc)
print(confusion_matrix(yc_test, pred_xgbc))
print(classification_report(yc_test, pred_xgbc))

XGBoost Classifier Accuracy: 0.9746557205260175
[[ 239]
 [ 458]]
precision recall f1-score support
0    0.96   0.98   0.97    239
1    0.98   0.96   0.97    219
accuracy          0.97   0.97   0.97    458
weighted avg    0.97   0.97   0.97    458

```



Accuracy Conclusion

Among the three classification models, XGBoost Classifier achieved the highest accuracy, demonstrating its strength in learning complex financial patterns. Decision Tree provided reasonable accuracy, while Logistic Regression served as a useful baseline but showed lower performance compared to tree-based models.

Conclusion & Future Scope

Conclusion :-

This project successfully demonstrated the use of machine learning models to predict gold prices and classify price trends. The Random Forest Regressor provided high accuracy for predicting actual gold prices, while among the classification models, XGBoost achieved the best performance. The results show that machine learning can effectively capture financial patterns and provide reliable predictions for gold price analysis.

Future Work :-

Future improvements may include adding more economic features like inflation, interest rates, and global market indicators to enhance accuracy. Advanced deep learning models such as LSTM can be used for time-series forecasting. The system can also be expanded into a real-time prediction dashboard using APIs or web applications. Additionally, hyperparameter tuning and ensemble stacking can further improve overall model performance.